

Assignment 1: Neural Machine Translation

Welcome to the first assignment of Course 4. Here, you will build an English-to-Portuguese neural machine translation (NMT) model using Long Short-Term Memory (LSTM) networks with attention. Machine translation is an important task in natural language processing and could be useful not only for translating one language to another but also for word sense disambiguation (e.g. determining whether the word "bank" refers to the financial bank, or the land alongside a river). Implementing this using just a Recurrent Neural Network (RNN) with LSTMs can work for short to medium length sentences but can result in vanishing gradients for very long sequences. To help with this, you will be adding an attention mechanism to allow the decoder to access all relevant parts of the input sentence regardless of its length. By completing this assignment, you will:

- Implement an encoder-decoder system with attention
- Build the NMT model from scratch using Tensorflow
- Generate translations using greedy and Minimum Bayes Risk (MBR) decoding

Table of Contents

- [1 - Data Preparation](#)
- [2 - NMT model with attention](#)
 - [Exercise 1 - Encoder](#)
 - [Exercise 2 - CrossAttention](#)
 - [Exercise 3 - Decoder](#)
 - [Exercise 4 - Translator](#)
- [3 - Training](#)
- [4 - Using the model for inference](#)
 - [Exercise 5 - translate](#)
- [5 - Minimum Bayes-Risk Decoding](#)
 - [Exercise 6 - rouge1_similarity](#)
 - [Exercise 7 - average_overlap](#)

```
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3' # Setting this env variable prevents TF warnings

import numpy as np
import tensorflow as tf
from collections import Counter
from utils import (sentences, train_data, val_data, english_vectorizer, portuguese_vectorizer)
```

1. Data Preparation

The text pre-processing bits have already been taken care of (if you are interested in this be sure to check the `utils.py` file). The steps performed can be summarized as:

- Reading the raw data from the text files
- Cleaning the data (using lowercase, adding space around punctuation, trimming whitespaces, etc)
- Splitting it into training and validation sets
- Adding the start-of-sentence and end-of-sentence tokens to every sentence
- Tokenizing the sentences
- Creating a Tensorflow dataset out of the tokenized sentences

Take a moment to inspect the raw sentences:

```
portuguese_sentences, english_sentences = sentences

print(f"English (to translate) sentence:\n\n{english_sentences[-5]}\n")
print(f"Portuguese (translation) sentence:\n\n{portuguese_sentences[-5]}")
```

English (to translate) sentence:

No matter how much you try to convince people that chocolate is vanilla, it'll still be chocolate, even though you may manage to convince yourself and a few others that it's vanilla.

Portuguese (translation) sentence:

Não importa o quanto você tenta convencer os outros de que chocolate é baunilha, ele ainda será chocolate, mesmo que você possa convencer a si mesmo e poucos outros de que é baunilha.

You don't have much use for the raw sentences so delete them to save memory:

```
del portuguese_sentences
del english_sentences
del sentences
```

Notice that you imported an `english_vectorizer` and a `portuguese_vectorizer` from `utils.py`. These were created using `tf.keras.layers.TextVectorization` and they provide interesting features such as ways to visualize the vocabulary and convert text into tokenized ids and vice versa. In fact, you can inspect the first ten words of the vocabularies for both languages:

```
print(f"First 10 words of the english vocabulary:\n\n{english_vectorizer.get_vocabulary()}")
print(f"First 10 words of the portuguese vocabulary:\n\n{portuguese_vectorizer.get_vocabulary()}")
```

First 10 words of the english vocabulary:

`['', '[UNK]', '[SOS]', '[EOS]', '.', 'tom', 'i', 'to', 'you', 'the']`

First 10 words of the portuguese vocabulary:

`['', '[UNK]', '[SOS]', '[EOS]', '.', 'tom', 'que', 'o', 'nao', 'eu']`

Notice that the first 4 words are reserved for special words. In order, these are:

- the empty string
- a special token to represent an unknown word
- a special token to represent the start of a sentence
- a special token to represent the end of a sentence

You can see how many words are in a vocabulary by using the `vocabulary_size` method:

```
# Size of the vocabulary
vocab_size_por = portuguese_vectorizer.vocabulary_size()
vocab_size_eng = english_vectorizer.vocabulary_size()

print(f"Portuguese vocabulary is made up of {vocab_size_por} words")
print(f"English vocabulary is made up of {vocab_size_eng} words")
```

Portuguese vocabulary is made up of 12000 words
 English vocabulary is made up of 12000 words

You can define `tf.keras.layers.StringLookup` objects that will help you map from words to ids and vice versa. Do this for the portuguese vocabulary since this will be useful later on when you decode the predictions from your model:

```
# This helps you convert from words to ids
word_to_id = tf.keras.layers.StringLookup(
    vocabulary=portuguese_vectorizer.get_vocabulary(),
    mask_token="",
    oov_token="[UNK]")
)

# This helps you convert from ids to words
id_to_word = tf.keras.layers.StringLookup(
    vocabulary=portuguese_vectorizer.get_vocabulary(),
    mask_token="",
    oov_token="[UNK]",
    invert=True,
```

Try it out for the special tokens and a random word:

```
unk_id = word_to_id("[UNK]")
sos_id = word_to_id("[SOS]")
eos_id = word_to_id("[EOS]")
baunilha_id = word_to_id("baunilha")

print(f"The id for the [UNK] token is {unk_id}")
print(f"The id for the [SOS] token is {sos_id}")
print(f"The id for the [EOS] token is {eos_id}")
print(f"The id for baunilha (vanilla) is {baunilha_id}")
```

```
The id for the [UNK] token is 1
The id for the [SOS] token is 2
The id for the [EOS] token is 3
The id for baunilha (vanilla) is 7079
```

Finally take a look at how the data that is going to be fed to the neural network looks like. Both `train_data` and `val_data` are of type `tf.data.Dataset` and are already arranged in batches of 64 examples. To get the first batch out of a `tf` dataset you can use the `take` method. To get the first example out of the batch you can slice the tensor and use the `numpy` method for nicer printing:

```
for (to_translate, sr_translation), translation in train_data.take(1):
    print(f"Tokenized english sentence:\n{to_translate[0, :].numpy()}\n")
    print(f"Tokenized portuguese sentence (shifted to the right):\n{sr_translation[0, :].numpy()}\n")
```

```
Tokenized english sentence:
[ 2 210 9 146 123 38 9 1672 4 3 0 0 0 0]
```

```
Tokenized portuguese sentence (shifted to the right):
[ 2 1085 7 128 11 389 37 2038 4 0 0 0 0 0]
```

```
Tokenized portuguese sentence:
[1085 7 128 11 389 37 2038 4 3 0 0 0 0 0]
```

There are a couple of important details to notice.

- Padding has already been applied to the tensors and the value used for this is 0
- Each example consists of 3 different tensors:
 - The sentence to translate

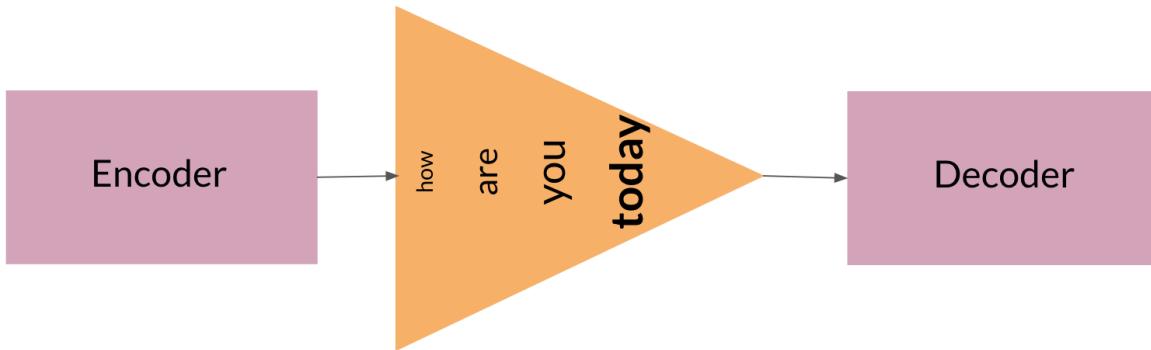
- The shifted-to-the-right translation
- The translation

The first two can be considered as the features, while the third one as the target. By doing this your model can perform Teacher Forcing as you saw in the lectures.

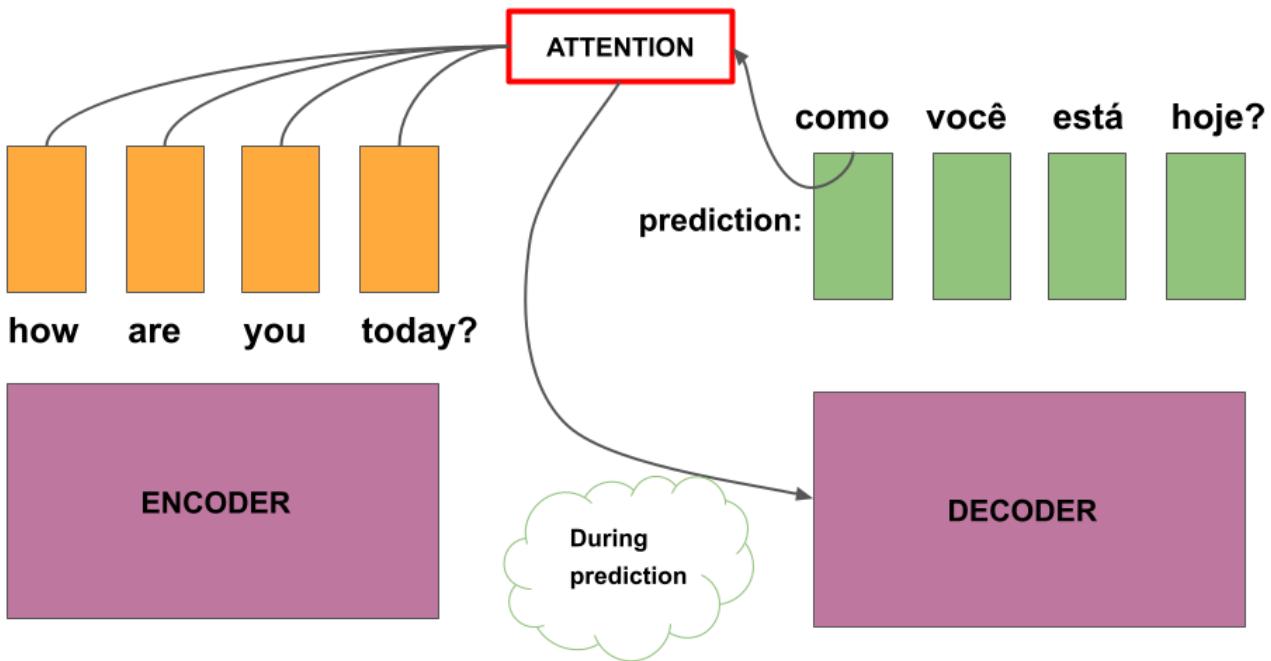
Now it is time to begin coding!

2. NMT model with attention

The model you will build uses an encoder-decoder architecture. This Recurrent Neural Network (RNN) takes in a tokenized version of a sentence in its encoder, then passes it on to the decoder for translation. As mentioned in the lectures, just using a regular sequence-to-sequence model with LSTMs will work effectively for short to medium sentences but will start to degrade for longer ones. You can picture it like the figure below where all of the context of the input sentence is compressed into one vector that is passed into the decoder block. You can see how this will be an issue for very long sentences (e.g. 100 tokens or more) because the context of the first parts of the input will have very little effect on the final vector passed to the decoder.



Adding an attention layer to this model avoids this problem by giving the decoder access to all parts of the input sentence. To illustrate, let's just use a 4-word input sentence as shown below. Remember that a hidden state is produced at each timestep of the encoder (represented by the orange rectangles). These are all passed to the attention layer and each are given a score given the current activation (i.e. hidden state) of the decoder. For instance, let's consider the figure below where the first prediction "como" is already made. To produce the next prediction, the attention layer will first receive all the encoder hidden states (i.e. orange rectangles) as well as the decoder hidden state when producing the word "como" (i.e. first green rectangle). Given this information, it will score each of the encoder hidden states to know which one the decoder should focus on to produce the next word. As a result of training, the model might have learned that it should align to the second encoder hidden state and subsequently assigns a high probability to the word "você". If we are using greedy decoding, we will output the said word as the next symbol, then restart the process to produce the next word until we reach an end-of-sentence prediction.



There are different ways to implement attention and the one we'll use for this assignment is the Scaled Dot Product Attention which has the form:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

You will dive deeper into this equation in the next week but for now, you can think of it as computing scores using queries (Q) and keys (K), followed by a multiplication of values (V) to get a context vector at a particular timestep of the decoder. This context vector is fed to the decoder RNN to get a set of probabilities for the next predicted word. The division by square root of the keys dimensionality ($\sqrt{d_k}$) is for improving

```
VOCAB_SIZE = 12000
UNITS = 256
```

Exercise 1 - Encoder

Your first exercise is to code the encoder part of the neural network. For this, complete the `Encoder` class below. Notice that in the constructor (the `__init__` method) you need to define all of the sublayers of the encoder and then use these sublayers during the forward pass (the `call` method).

The encoder consists of the following layers:

- [Embedding](#). For this layer you need to define the appropriate `input_dim` and `output_dim` and let it know that you are using '0' as padding, which can be done by using the appropriate value for the `mask_zero` parameter.
- [Bidirectional LSTM](#). In TF you can implement bidirectional behaviour for RNN-like layers. This part is already taken care of but you will need to specify the appropriate type of layer as well as its parameters. In particular you need to set the appropriate number of units and make sure that the LSTM returns the full sequence and not only the last output, which can be done by using the appropriate value for the `return_sequences` parameter.

You need to define the forward pass using the syntax of TF's [functional API](#). What this means is that you chain function calls together to define your network like this:

```
encoder_input = keras.Input(shape=(28, 28, 1), name="original_img")
x = layers.Conv2D(16, 3, activation="relu")(encoder_input)
x = layers.MaxPooling2D(3)(x)
```

```
# GRADED CLASS: Encoder
class Encoder(tf.keras.layers.Layer):
    def __init__(self, vocab_size, units):
        """Initializes an instance of this class

    Args:
        vocab_size (int): Size of the vocabulary
        units (int): Number of units in the LSTM layer
    """
    super(Encoder, self).__init__()

    ### START CODE HERE ###

    # Embedding Layer: converts token IDs to dense vectors
    self.embedding = tf.keras.layers.Embedding(
        input_dim=vocab_size,
        output_dim=units,
        mask_zero=True # ignore padding (0s)
    )

    # Bidirectional LSTM: processes sequence in both directions
    self.rnn = tf.keras.layers.Bidirectional(
        tf.keras.layers.LSTM(
            units=units,
            return_sequences=True # return full sequence (not just final state)
        ),
        merge_mode="sum" # sum forward and backward outputs
    )

    ### END CODE HERE ###

    def call(self, context):
        """Forward pass of this layer

    Args:
        context (tf.Tensor): The sentence to translate

    Returns:
        tf.Tensor: Encoded sentence to translate
    """

    ### START CODE HERE ###

    # Step 1: Embed the input tokens
    x = self.embedding(context)

    # Step 2: Process through the BiLSTM
    x = self.rnn(x)

    ### END CODE HERE ###
```

```
# Do a quick check of your implementation

# Create an instance of your class
encoder = Encoder(VOCAB_SIZE, UNITS)

# Pass a batch of sentences to translate from english to portuguese
encoder_output = encoder(to_translate)

print(f'Tensor of sentences in english has shape: {to_translate.shape}\n')
print(f'Encoder output has shape: {encoder_output.shape}\n')
```

Tensor of sentences in english has shape: (64, 14)

Encoder output has shape: (64, 14, 256)

Expected Output

Tensor of sentences in english has shape: (64, 14)

Encoder output has shape: (64, 14, 256)

```
# Test your code!
```

~~w1_units.assert_equal(encoder_output, Encoder)~~

All tests passed!

Exercise 2 - CrossAttention

Your next exercise is to code the layer that will perform cross attention between the original sentences and the translations. For this, complete the `CrossAttention` class below. Notice that in the constructor (the `__init__` method) you need to define all of the sublayers and then use these sublayers during the forward pass (the `call` method). For this particular case some of these bits are already taken care of.

The cross attention consists of the following layers:

- [MultiHeadAttention](#). For this layer you need to define the appropriate `key_dim`, which is the size of the key and query tensors. You will also need to set the number of heads to 1 since you aren't implementing multi head attention but attention between two tensors. The reason why this layer is preferred over [Attention](#) is that it allows simpler code during the forward pass.

A couple of things to notice:

- You need a way to pass both the output of the attention alongside the shifted-to-the-right translation (since this cross attention happens in the decoder side). For this you will use an [Add](#) layer so that the original dimension is preserved, which would not happen if you use something like a [Concatenate](#) layer.
- Layer normalization is also performed for better stability of the network by using a [LayerNormalization](#) layer.
- You don't need to worry about these last steps as these are already solved.

```

# GRADED CLASS: CrossAttention
class CrossAttention(tf.keras.layers.Layer):
    def __init__(self, units):
        """Initializes an instance of this class

    Args:
        units (int): Number of units in the LSTM layer
    """
    super().__init__()

    ### START CODE HERE ###

    # Single-head attention layer (scaled dot-product attention)
    self.mha = tf.keras.layers.MultiHeadAttention(
        key_dim=units, # size of key/query vectors
        num_heads=1     # single attention head
    )

    ### END CODE HERE ###

    # Normalization and residual connection
    self.layernorm = tf.keras.layers.LayerNormalization()
    self.add = tf.keras.layers.Add()

    def call(self, context, target):
        """Forward pass of this layer

    Args:
        context (tf.Tensor): Encoded sentence to translate
        target (tf.Tensor): The embedded shifted-to-the-right translation

    Returns:
        tf.Tensor: Cross attention between context and target
    """
    ### START CODE HERE ###

    # Perform attention: query = target (decoder input), value = context (encoder out
    attn_output = self.mha(
        query=target,
        value=context,
        key=context
    )

    ### END CODE HERE ###

    # Add residual connection and normalize
    x = self.add([target, attn_output])
    x = self.layernorm(x)

    return x

```

```
# Do a quick check of your implementation

# Create an instance of your class
attention_layer = CrossAttention(UNITS)

# The attention layer expects the embedded sr-translation and the context
# The context (encoder_output) is already embedded so you need to do this for sr_translat
sr_translation_embed = tf.keras.layers.Embedding(VOCAB_SIZE, output_dim=UNITS, mask_zero=True)

# Compute the cross attention
attention_result = attention_layer(encoder_output, sr_translation_embed)

print(f'Tensor of contexts has shape: {encoder_output.shape}')
print(f'Tensor of translations has shape: {sr_translation_embed.shape}')
print(f'Tensor of attention scores has shape: {attention_result.shape}')


Tensor of contexts has shape: (64, 14, 256)
Tensor of translations has shape: (64, 15, 256)
Tensor of attention scores has shape: (64, 15, 256)
```

Expected Output

```
Tensor of contexts has shape: (64, 14, 256)
Tensor of translations has shape: (64, 15, 256)
Tensor of attention scores has shape: (64, 15, 256)
```

```
# Test your code!
```

```
.../unit-test/test_cross_attention(CrossAttention)
All tests passed!
```

Exercise 3 - Decoder

Now you will implement the decoder part of the neural network by completing the `Decoder` class below. Notice that in the constructor (the `__init__` method) you need to define all of the sublayers of the decoder and then use these sublayers during the forward pass (the `call` method).

The decoder consists of the following layers:

- Embedding. For this layer you need to define the appropriate `input_dim` and `output_dim` and let it know that you are using '0' as padding, which can be done by using the appropriate value for the `mask_zero` parameter.
- Pre-attention LSTM. Unlike in the encoder in which you used a Bidirectional LSTM, here you will use a vanilla LSTM. Don't forget to set the appropriate number of units and make sure that the LSTM returns the full sequence and not only the last output, which can be done by using the appropriate value for the `return_sequences` parameter. It is very important that this layer returns the state since this will be needed for inference so make sure to set the `return_state` parameter accordingly. Notice that LSTM layers return state as a tuple of two tensors called `memory_state` and `carry_state`, **however these names have been changed to better reflect what you have seen in the lectures to `hidden_state` and `cell_state` respectively**.
- The attention layer that performs cross attention between the sentence to translate and the right-shifted translation. Here you need to use the `CrossAttention` layer you defined in the previous exercise.
- Post-attention LSTM. Another LSTM layer. For this one you don't need it to return the state.
- Finally a Dense layer. This one should have the same number of units as the size of the vocabulary since you expect it to compute the logits for every possible word in the vocabulary. Make sure to use

a `logsoftmax` activation function for this one, which you can get as [`tf.nn.log_softmax`](#).


```

# GRADED CLASS: Decoder
class Decoder(tf.keras.layers.Layer):
    def __init__(self, vocab_size, units):
        """Initializes an instance of this class

    Args:
        vocab_size (int): Size of the vocabulary
        units (int): Number of units in the LSTM layer
    """
        super(Decoder, self).__init__()

    ### START CODE HERE ###

    # The embedding layer
    self.embedding = tf.keras.layers.Embedding(
        input_dim=vocab_size,
        output_dim=units,
        mask_zero=True
    )

    # The RNN before attention
    self.pre_attention_rnn = tf.keras.layers.LSTM(
        units=units,
        return_sequences=True,
        return_state=True
    )

    # The attention layer (use the CrossAttention class you implemented)
    self.attention = CrossAttention(units)

    # The RNN after attention
    self.post_attention_rnn = tf.keras.layers.LSTM(
        units=units,
        return_sequences=True
    )

    # The dense layer with Logsoftmax activation
    self.output_layer = tf.keras.layers.Dense(
        units=vocab_size,
        activation=tf.nn.log_softmax
    )

    ### END CODE HERE ###

    def call(self, context, target, state=None, return_state=False):
        """Forward pass of this layer

    Args:
        context (tf.Tensor): Encoded sentence to translate
        target (tf.Tensor): The shifted-to-the-right translation
        state (list[tf.Tensor, tf.Tensor], optional): Hidden state of the pre-attention
        return_state (bool, optional): If set to true return the hidden states of the
    """

    Returns:
        tf.Tensor: The log_softmax probabilities of predicting a particular token
    """
    ### START CODE HERE ###

    # Step 1: Embed the input
    x = self.embedding(target)

```

```

# Step 2: Pass through the pre-attention LSTM
# This returns output + two states (hidden, cell)
x, hidden_state, cell_state = self.pre_attention_rnn(x, initial_state=state)

# Step 3: Apply cross-attention between encoder context and decoder pre-attention
x = self.attention(context, x)

# Step 4: Pass through the post-attention LSTM
x = self.post_attention_rnn(x)

# Step 5: Compute log-softmax logits over the vocabulary
logits = self.output_layer(x)

### END CODE HERE ###

if return_state:
    return logits, [hidden_state, cell_state]

return logits

```

Do a quick check of your implementation

```

# Create an instance of your class
decoder = Decoder(VOCAB_SIZE, UNITS)

# Notice that you don't need the embedded version of sr_translation since this is done in
logits = decoder(encoder_output, sr_translation)

print(f'Tensor of contexts has shape: {encoder_output.shape}')
print(f'Tensor of right-shifted translations has shape: {sr_translation.shape}')
print(f'Tensor of logits has shape: {logits.shape}')

Tensor of contexts has shape: (64, 14, 256)
Tensor of right-shifted translations has shape: (64, 15)
Tensor of logits has shape: (64, 15, 12000)

```

Expected Output

```

Tensor of contexts has shape: (64, 14, 256)
Tensor of right-shifted translations has shape: (64, 15)
Tensor of logits has shape: (64, 15, 12000)

```

Test your code!

All tests passed!

Exercise 4 - Translator

Now you have to put together all of the layers you previously coded into an actual model. For this, complete the `Translator` class below. Notice how unlike the Encoder and Decoder classes inherited from `tf.keras.layers.Layer`, the `Translator` class inherits from `tf.keras.Model`.

Remember that `train_data` will yield a tuple with the sentence to translate and the shifted-to-the-right translation, which are the "features" of the model. This means that the inputs of your network will be tuples containing context and targets.

```

# GRADED CLASS: Translator
class Translator(tf.keras.Model):
    def __init__(self, vocab_size, units):
        """Initializes an instance of this class

        Args:
            vocab_size (int): Size of the vocabulary
            units (int): Number of units in the LSTM layer
        """
        super().__init__()

    ### START CODE HERE ###

    # Define the encoder
    self.encoder = Encoder(vocab_size, units)

    # Define the decoder
    self.decoder = Decoder(vocab_size, units)

    ### END CODE HERE ###

    def call(self, inputs):
        """Forward pass of this layer

        Args:
            inputs (tuple(tf.Tensor, tf.Tensor)): Tuple containing the context (sentence
                                                and the target (shifted-to-the-right tr

        Returns:
            tf.Tensor: The log_softmax probabilities of predicting a particular token
        """
    ### START CODE HERE ###

    # Unpack the inputs tuple
    context, target = inputs

    # Encode the context (English sentence)
    encoded_context = self.encoder(context)

    # Decode to generate the logits (Portuguese translation)
    logits = self.decoder(encoded_context, target)

    ### END CODE HERE ###

```

...
...

```

# Do a quick check of your implementation

# Create an instance of your class
translator = Translator(VOCAB_SIZE, UNITS)

# Compute the Logits for every word in the vocabulary
logits = translator((to_translate, sr_translation))

print(f'Tensor of sentences to translate has shape: {to_translate.shape}')
print(f'Tensor of right-shifted translations has shape: {sr_translation.shape}')
print(f'Tensor of logits has shape: {logits.shape}')

Tensor of sentences to translate has shape: (64, 14)
Tensor of right-shifted translations has shape: (64, 15)
Tensor of logits has shape: (64, 15, 12000)

```

Expected Output

```
Tensor of sentences to translate has shape: (64, 14)
Tensor of right-shifted translations has shape: (64, 15)
Tensor of logits has shape: (64, 15, 12000)
```

```
.../nmt_with_attention/Translation_Encoder_Decoder
```

```
All tests passed!
```

3. Training

Now that you have an untrained instance of the NMT model, it is time to train it. You can use the `compile_and_train` function below to achieve this:

```
def compile_and_train(model, epochs=20, steps_per_epoch=500):
    model.compile(optimizer="adam", loss=masked_loss, metrics=[masked_acc, masked_loss])

    history = model.fit(
        train_data.repeat(),
        epochs=epochs,
        steps_per_epoch=steps_per_epoch,
        validation_data=val_data,
        validation_steps=50,
        callbacks=[tf.keras.callbacks.EarlyStopping(patience=3)],
    )

    return model, history
```

```
# Train the translator (this takes some minutes so feel free to take a break)
```

```
!trained_translator.history = compile_and_train(translator)
```

Epoch 1/20

500/500 [=====] - 47s 67ms/step - loss: 5.1603 -
masked_acc: 0.2155 - masked_loss: 5.1628 - val_loss: 4.3576 - val_masked_a
cc: 0.3170 - val_masked_loss: 4.3587

Epoch 2/20

500/500 [=====] - 17s 33ms/step - loss: 3.7966 -
masked_acc: 0.4053 - masked_loss: 3.7973 - val_loss: 3.1185 - val_masked_a
cc: 0.4894 - val_masked_loss: 3.1203

Epoch 3/20

500/500 [=====] - 16s 32ms/step - loss: 2.7888 -
masked_acc: 0.5391 - masked_loss: 2.7903 - val_loss: 2.4437 - val_masked_a
cc: 0.5780 - val_masked_loss: 2.4442

Epoch 4/20

500/500 [=====] - 16s 32ms/step - loss: 2.2626 -
masked_acc: 0.6129 - masked_loss: 2.2634 - val_loss: 2.0084 - val_masked_a
cc: 0.6439 - val_masked_loss: 2.0087

Epoch 5/20

500/500 [=====] - 15s 31ms/step - loss: 1.9037 -
masked_acc: 0.6640 - masked_loss: 1.9048 - val_loss: 1.7488 - val_masked_a
cc: 0.6812 - val_masked_loss: 1.7492

Epoch 6/20

500/500 [=====] - 15s 30ms/step - loss: 1.6415 -
masked_acc: 0.6965 - masked_loss: 1.6425 - val_loss: 1.6149 - val_masked_a
cc: 0.7016 - val_masked_loss: 1.6144

Epoch 7/20

500/500 [=====] - 15s 31ms/step - loss: 1.5241 -
masked_acc: 0.7133 - masked_loss: 1.5248 - val_loss: 1.5015 - val_masked_a
cc: 0.7114 - val_masked_loss: 1.5019

Epoch 8/20

500/500 [=====] - 16s 32ms/step - loss: 1.4169 -
masked_acc: 0.7275 - masked_loss: 1.4178 - val_loss: 1.4057 - val_masked_a
cc: 0.7248 - val_masked_loss: 1.4075

Epoch 9/20

500/500 [=====] - 16s 31ms/step - loss: 1.3380 -
masked_acc: 0.7374 - masked_loss: 1.3393 - val_loss: 1.3131 - val_masked_a
cc: 0.7376 - val_masked_loss: 1.3133

Epoch 10/20

500/500 [=====] - 16s 31ms/step - loss: 1.2089 -
masked_acc: 0.7526 - masked_loss: 1.2102 - val_loss: 1.2516 - val_masked_a
cc: 0.7478 - val_masked_loss: 1.2525

Epoch 11/20

500/500 [=====] - 15s 31ms/step - loss: 1.1022 -
masked_acc: 0.7651 - masked_loss: 1.1031 - val_loss: 1.1885 - val_masked_a
cc: 0.7604 - val_masked_loss: 1.1888

Epoch 12/20

500/500 [=====] - 15s 30ms/step - loss: 1.0790 -
masked_acc: 0.7700 - masked_loss: 1.0799 - val_loss: 1.1513 - val_masked_a
cc: 0.7588 - val_masked_loss: 1.1515

Epoch 13/20

500/500 [=====] - 15s 31ms/step - loss: 1.0491 -
masked_acc: 0.7729 - masked_loss: 1.0503 - val_loss: 1.1381 - val_masked_a
cc: 0.7609 - val_masked_loss: 1.1390

Epoch 14/20

500/500 [=====] - 15s 31ms/step - loss: 1.0251 -
masked_acc: 0.7765 - masked_loss: 1.0262 - val_loss: 1.1032 - val_masked_a
cc: 0.7685 - val_masked_loss: 1.1035

Epoch 15/20

500/500 [=====] - 16s 31ms/step - loss: 0.9190 -
masked_acc: 0.7911 - masked_loss: 0.9201 - val_loss: 1.0783 - val_masked_a
cc: 0.7698 - val_masked_loss: 1.0800

Epoch 16/20

```
500/500 [=====] - 15s 31ms/step - loss: 0.8697 -  
masked_acc: 0.7978 - masked_loss: 0.8703 - val_loss: 1.0749 - val_masked_a  
cc: 0.7714 - val_masked_loss: 1.0765  
Epoch 17/20  
500/500 [=====] - 15s 30ms/step - loss: 0.8760 -  
masked_acc: 0.7950 - masked_loss: 0.8767 - val_loss: 1.0194 - val_masked_a  
cc: 0.7770 - val_masked_loss: 1.0207  
Epoch 18/20  
500/500 [=====] - 15s 31ms/step - loss: 0.8751 -  
masked_acc: 0.7958 - masked_loss: 0.8758 - val_loss: 1.0120 - val_masked_a  
cc: 0.7809 - val_masked_loss: 1.0130  
Epoch 19/20  
500/500 [=====] - 15s 30ms/step - loss: 0.8540 -  
masked_acc: 0.7993 - masked_loss: 0.8547 - val_loss: 0.9915 - val_masked_a  
cc: 0.7819 - val_masked_loss: 0.9913  
Epoch 20/20  
500/500 [=====] - 15s 31ms/step - loss: 0.7479 -  
masked_acc: 0.8161 - masked_loss: 0.7486 - val_loss: 0.9973 - val_masked_a  
cc: 0.7831 - val_masked_loss: 0.9980
```

4. Using the model for inference

Now that your model is trained you can use it for inference. To help you with this the `generate_next_token` function is provided. Notice that this function is meant to be used inside a for-loop, so you feed to it the information of the previous step to generate the information of the next step. In particular you need to keep track of the state of the pre-attention LSTM in the decoder and if you are done with the translation. Also notice that a `temperature` variable is introduced which determines how to select the next token given the predicted logits:

```
def generate_next_token(decoder, context, next_token, done, state, temperature=0.0):
    """Generates the next token in the sequence

Args:
    decoder (Decoder): The decoder
    context (tf.Tensor): Encoded sentence to translate
    next_token (tf.Tensor): The predicted next token
    done (bool): True if the translation is complete
    state (list[tf.Tensor, tf.Tensor]): Hidden states of the pre-attention LSTM layer
    temperature (float, optional): The temperature that controls the randomness of the
                                    selection of the next token

Returns:
    tuple(tf.Tensor, np.float, list[tf.Tensor, tf.Tensor], bool): The next token, log
    ..."""

# Get the logits and state from the decoder
logits, state = decoder(context, next_token, state=state, return_state=True)

# Trim the intermediate dimension
logits = logits[:, -1, :]

# If temp is 0 then next_token is the argmax of logits
if temperature == 0.0:
    next_token = tf.argmax(logits, axis=-1)

# If temp is not 0 then next_token is sampled out of logits
else:
    logits = logits / temperature
    next_token = tf.random.categorical(logits, num_samples=1)

# Trim dimensions of size 1
logits = tf.squeeze(logits)
next_token = tf.squeeze(next_token)

# Get the logit of the selected next_token
logit = logits[next_token].numpy()

# Reshape to (1,1) since this is the expected shape for text encoded as TF tensors
next_token = tf.reshape(next_token, shape=(1,1))

# If next_token is End-of-Sentence token you are done
if next_token == eos_id:
    done = True
```

See how it works by running the following cell:

```

# PROCESS SENTENCE TO TRANSLATE AND ENCODE

# A sentence you wish to translate
eng_sentence = "I love languages"

# Convert it to a tensor
texts = tf.convert_to_tensor(eng_sentence)[tf.newaxis]

# Vectorize it and pass it through the encoder
context = english_vectorizer(texts).to_tensor()
context = encoder(context)

# SET STATE OF THE DECODER

# Next token is Start-of-Sentence since you are starting fresh
next_token = tf.fill((1,1), sos_id)

# Hidden and Cell states of the LSTM can be mocked using uniform samples
state = [tf.random.uniform((1, UNITS)), tf.random.uniform((1, UNITS))]

# You are not done until next token is EOS token
done = False

# Generate next token
next_token, logit, state, done = generate_next_token(decoder, context, next_token, done,
                                                       **{"Next token": next_token, "Logit": logit, "Done": done})

Next token: [[6188]]
Logit: -18.7976
Done? False

```

Exercise 5 - translate

Now you can put everything together to translate a given sentence. For this, complete the `translate` function below. This function will take care of the following steps:

- Process the sentence to translate and encode it
- Set the initial state of the decoder
- Get predictions of the next token (starting with the <SOS> token) for a maximum of iterations (in case the <EOS> token is never returned)
- Return the translated text (as a string), the logit of the last iteration (this helps measure how certain was that the sequence was translated in its totality) and the translation in token format.

Hints:

- The previous cell provides a lot of insights on how this function should work, so if you get stuck refer to it.
- Some useful docs:
 - [tf.newaxis](#)
 - [tf.fill](#)
 - [tf.zeros](#)

IMPORTANT NOTE: Due to randomness processes involving tensorflow training and weight initializing, the results below may vary a lot, even if you retrain your model in the same session.


```

# GRADED FUNCTION: translate
def translate(model, text, max_length=50, temperature=0.0):
    """Translate a given sentence from English to Portuguese

Args:
    model (tf.keras.Model): The trained translator
    text (string): The sentence to translate
    max_length (int, optional): The maximum length of the translation. Defaults to 50
    temperature (float, optional): The temperature that controls the randomness of the
                                    generated tokens

Returns:
    tuple(str, np.float, tf.Tensor): The translation, logit that predicted <EOS> token
"""

# Lists to save tokens and logits
tokens, logits = [], []

### START CODE HERE ###

# PROCESS THE SENTENCE TO TRANSLATE

# Convert the original string into a tensor
text = tf.convert_to_tensor([text])

# Vectorize the text using the English vectorizer
context = english_vectorizer(text).to_tensor()

# Get the encoded context (pass the context through the encoder)
context = model.encoder(context)

# INITIAL STATE OF THE DECODER

# First token should be SOS token with shape (1,1)
next_token = tf.fill((1, 1), sos_id)

# Initial hidden and cell states should be tensors of zeros with shape (1, UNITS)
state = [tf.zeros((1, UNITS)), tf.zeros((1, UNITS))]

# You are done when you draw an EOS token as next token (initial state is False)
done = False

# Iterate for max_length iterations
for _ in range(max_length):
    # Generate the next token
    try:
        next_token, logit, state, done = generate_next_token(
            decoder=model.decoder,
            context=context,
            next_token=next_token,
            done=done,
            state=state,
            temperature=temperature
        )
    except Exception as e:
        raise Exception("Problem generating the next token") from e

    # If done then break out of the loop
    if done:
        break

    # Add next_token to the list of tokens
    tokens.append(next_token)

```

```

# Add Logit to the list of logits
logits.append(logit)

### END CODE HERE ###

# Concatenate all tokens into a tensor
tokens = tf.concat(tokens, axis=-1)

# Convert the translated tokens into text
translation = tf.squeeze(tokens_to_text(tokens, id_to_word))
translation = translation.numpy().decode()

return translation, logits, tokens

```

Try your function with temperature of 0, which will yield a deterministic output and is equivalent to a greedy decoding:

```

# Running this cell multiple times should return the same output since temp is 0

temp = 0.0
original_sentence = "I love languages"

translation, logit, tokens = translate(trained_translator, original_sentence, temperature)

print(f"Temperature: {temp}\nOriginal sentence: {original_sentence}\nTranslation: {translati
Temperature: 0.0

Original sentence: I love languages
Translation: eu adoro idiomas para idiomas .
Translation tokens:[[ 9 564 850 18 850 4]]
Logit: -0.413

```

Try your function with temperature of 0.7 (stochastic output):

```

# Running this cell multiple times should return different outputs since temp is not 0
# You can try different temperatures

temp = 0.7
original_sentence = "I love languages"

translation, logit, tokens = translate(trained_translator, original_sentence, temperature)

print(f"Temperature: {temp}\nOriginal sentence: {original_sentence}\nTranslation: {translati
Temperature: 0.7

Original sentence: I love languages
Translation: eu amo idiomas a vida .
Translation tokens:[[ 9 522 850 12 195 4]]
Logit: -2.045

```

```
u1.unittest.test_translate(translate, trained_translator)
All tests passed!
```

5. Minimum Bayes-Risk Decoding

As mentioned in the lectures, getting the most probable token at each step may not necessarily produce the best results. Another approach is to do Minimum Bayes Risk Decoding or MBR. The general steps to implement this are:

- Take several random samples
- Score each sample against all other samples
- Select the one with the highest score

You will be building helper functions for these steps in the following sections.

With the ability to generate different translations by setting different temperature values you can do what you saw in the lectures and generate a bunch of translations and then determine which one is the best

```
def generate_samples(model, text, n_samples=4, temperature=0.6):
    samples, log_probs = [], []
    # Iterate for n_samples iterations
    for _ in range(n_samples):
        # Save the Logit and the translated tensor
        _, logp, sample = translate(model, text, temperature=temperature)

        # Save the translated tensors
        samples.append(np.squeeze(sample.numpy()).tolist())

        # Save the logits
        log_probs.append(logp)
    return samples, log_probs
```

```
samples, log_probs = generate_samples(trained_translator, 'I love languages')
```

```
for s, l in zip(samples, log_probs):
    print(f"Translated tensor: {s} has logit: {l:.5f}")
```

```
Translated tensor: [9, 564, 850, 18, 7, 811, 4] has logit: -0.584
Translated tensor: [9, 522, 850, 18, 850, 4] has logit: -0.985
Translated tensor: [9, 522, 850, 7, 210, 4] has logit: -2.901
Translated tensor: [9, 564, 850, 67, 811, 4] has logit: -0.676
```

Comparing overlaps

Now that you can generate multiple translations it is time to come up with a method to measure the goodness of each one. As you saw in the lectures, one way to achieve this is by comparing each sample against the others.

There are several metrics you can use for this purpose, as shown in the lectures and you can try experimenting with any one of these. For this assignment, you will be calculating scores for **unigram overlaps**.

One of these metrics is the widely used yet simple [Jaccard similarity](#) which gets the intersection over union of two sets. The `jaccard_similarity` function returns this metric for any pair of candidate and reference translations:

```
def jaccard_similarity(candidate, reference):

    # Convert the lists to sets to get the unique tokens
    candidate_set = set(candidate)
    reference_set = set(reference)

    # Get the set of tokens common to both candidate and reference
    common_tokens = candidate_set.intersection(reference_set)

    # Get the set of all tokens found in either candidate or reference
    all_tokens = candidate_set.union(reference_set)

    # Compute the percentage of overlap (divide the number of common tokens by the number
    overlap = len(common_tokens) / len(all_tokens)

    return overlap
```

```
l1 = [1, 2, 3]
l2 = [1, 2, 3, 4]

js = jaccard_similarity(l1, l2)
```

jaccard similarity between lists: [1, 2, 3] and [1, 2, 3, 4] is 0.750

Expected Output

jaccard similarity between tensors: [1, 2, 3] and [1, 2, 3, 4] is 0.750

Exercise 6 - rouge1_similarity

Jaccard similarity is good but a more commonly used metric in machine translation is the ROUGE score. For unigrams, this is called ROUGE-1 and as shown in the lectures, you can output the scores for both precision and recall when comparing two samples. To get the final score, you will want to compute the F1-score as given by:

$$score = 2 * \frac{(precision * recall)}{(precision + recall)}$$

For the implementation of the `rouge1_similarity` function you want to use the [Counter](#) class from the Python standard library:

```
# GRADED FUNCTION: rouge1_similarity
def rouge1_similarity(candidate, reference):
    """Computes the ROUGE 1 score between two token lists

    Args:
        candidate (list[int]): Tokenized candidate translation
        reference (list[int]): Tokenized reference translation

    Returns:
        float: Overlap between the two token lists
    """
    ### START CODE HERE ###

    # Make frequency tables for candidate and reference
    candidate_word_counts = Counter(candidate)
    reference_word_counts = Counter(reference)

    # Initialize overlap
    overlap = 0.0

    # Iterate over tokens in candidate
    for token in candidate_word_counts.keys():
        # Count of current token in candidate and reference
        token_count_candidate = candidate_word_counts[token]
        token_count_reference = reference_word_counts[token]

        # Add minimum count to overlap (handles duplicates)
        overlap += min(token_count_candidate, token_count_reference)

    # Compute precision and recall
    precision = overlap / len(candidate) if len(candidate) > 0 else 0
    recall = overlap / len(reference) if len(reference) > 0 else 0

    # Compute F1 score (ROUGE-1)
    if precision + recall != 0:
        f1_score = 2 * (precision * recall) / (precision + recall)
        return f1_score

    ### END CODE HERE ###
```

```
l1 = [1, 2, 3]
l2 = [1, 2, 3, 4]
```

```
r1s = rouge1_similarity(l1, l2)
```

```
rouge 1 similarity between lists: [1, 2, 3] and [1, 2, 3, 4] is 0.857
```

Expected Output

```
rouge 1 similarity between lists: [1, 2, 3] and [1, 2, 3, 4] is 0.857
```

```
u1.unittest.test_rouge1_similarity(rouge1_similarity)
```

All tests passed!

Computing the Overall Score

You will now build a function to generate the overall score for a particular sample. As mentioned in the lectures, you need to compare each sample with all other samples. For instance, if we generated 30 sentences, we will need to compare sentence 1 to sentences 2 through 30. Then, we compare sentence 2 to sentences 1 and 3 through 30, and so forth. At each step, we get the average score of all comparisons to get the overall score for a particular sample. To illustrate, these will be the steps to generate the scores of a 4-sample list.

- Get similarity score between sample 1 and sample 2
- Get similarity score between sample 1 and sample 3
- Get similarity score between sample 1 and sample 4
- Get average score of the first 3 steps. This will be the overall score of sample 1
- Iterate and repeat until samples 1 to 4 have overall scores.

The results will be stored in a dictionary for easy lookups.

Exercise 7 - average_overlap

Complete the `average_overlap` function below which should implement the process described above:

```

# GRADED FUNCTION: average_overlap
def average_overlap(samples, similarity_fn):
    """Computes the arithmetic mean of each candidate sentence in the samples

Args:
    samples (list[list[int]]): Tokenized version of translated sentences
    similarity_fn (Function): Similarity function used to compute the overlap

Returns:
    dict[int, float]: A dictionary mapping the index of each translation to its score
"""

# Initialize dictionary
scores = {}

# Iterate through all samples (enumerate helps keep track of indexes)
for index_candidate, candidate in enumerate(samples):

    ### START CODE HERE ###

    # Initially overlap is zero
    overlap = 0.0

    # Iterate through all samples
    for index_sample, sample in enumerate(samples):

        # Skip if the candidate index is the same as the sample index
        if index_candidate == index_sample:
            continue

        # Get the overlap between candidate and sample using the similarity function
        sample_overlap = similarity_fn(candidate, sample)

        # Add the sample overlap to the total overlap
        overlap += sample_overlap

    ### END CODE HERE ###

    # Get the score for the candidate by computing the average
    score = overlap / (len(samples) - 1)

    # Only use 3 decimal points
    score = round(score, 3)

    # Save the score in the dictionary using index as the key
    scores[index_candidate] = score

```

Test with Jaccard similarity

```
l1 = [1, 2, 3]
l2 = [1, 2, 4]
l3 = [1, 2, 4, 5]

avg_ovlp = average_overlap([l1, l2, l3], jaccard_similarity)
```

average overlap between lists: [1, 2, 3], [1, 2, 4] and [1, 2, 4, 5] using Jaccard similarity is:

```
{0: 0.45, 1: 0.625, 2: 0.575}
```

Expected Output

average overlap between lists: [1, 2, 3], [1, 2, 4] and [1, 2, 4, 5] using Jaccard similarity is:

```
{0: 0.45, 1: 0.625, 2: 0.575}
```

Test with Rouge1 similarity

```
l1 = [1, 2, 3]
l2 = [1, 4]
l3 = [1, 2, 4, 5]
l4 = [5,6]

avg_ovlp = average_overlap([l1, l2, l3, l4], rouge1_similarity)
```

average overlap between lists: [1, 2, 3], [1, 4], [1, 2, 4, 5] and [5, 6] using Rouge1 similarity is:

```
{0: 0.324, 1: 0.356, 2: 0.524, 3: 0.111}
```

Expected Output

average overlap between lists: [1, 2, 3], [1, 4], [1, 2, 4, 5] and [5, 6] using Rouge1 similarity is:

```
{0: 0.324, 1: 0.356, 2: 0.524, 3: 0.111}
```

`l1 = [1, 2, 3]
l2 = [1, 4]
l3 = [1, 2, 4, 5]
l4 = [5,6]`

All tests passed!

In practice, it is also common to see the weighted mean being used to calculate the overall score instead of just the arithmetic mean. This is implemented in the `weighted_avg_overlap` function below and you can use it in your experiments to see which one will give better results:

```

def weighted_avg_overlap(samples, log_probs, similarity_fn):

    # Scores dictionary
    scores = {}

    # Iterate over the samples
    for index_candidate, candidate in enumerate(samples):

        # Initialize overlap and weighted sum
        overlap, weight_sum = 0.0, 0.0

        # Iterate over all samples and Log probabilities
        for index_sample, (sample, logp) in enumerate(zip(samples, log_probs)):

            # Skip if the candidate index is the same as the sample index
            if index_candidate == index_sample:
                continue

            # Convert Log probability to linear scale
            sample_p = float(np.exp(logp))

            # Update the weighted sum
            weight_sum += sample_p

            # Get the unigram overlap between candidate and sample
            sample_overlap = similarity_fn(candidate, sample)

            # Update the overlap
            overlap += sample_p * sample_overlap

        # Compute the score for the candidate
        score = overlap / weight_sum

        # Only use 3 decimal points
        score = round(score, 3)

        # Save the score in the dictionary. use index as the key.
        scores[index_candidate] = score
    
```

```

l1 = [1, 2, 3]
l2 = [1, 2, 4]
l3 = [1, 2, 4, 5]
log_probs = [0.4, 0.2, 0.5]

w_avg_ovlp = weighted_avg_overlap([l1, l2, l3], log_probs, jaccard_similarity)

```

weighted average overlap using Jaccard similarity is:

```
{0: 0.443, 1: 0.631, 2: 0.558}
```

mbr_decode

You will now put everything together in the the `mbr_decode` function below. This final step is not graded as this function is just a wrapper around all the cool stuff you have coded so far!

You can use it to play around, trying different numbers of samples, temperatures and similarity functions!

```

def mbr_decode(model, text, n_samples=5, temperature=0.6, similarity_fn=jaccard_similarit

    # Generate samples
    samples, log_probs = generate_samples(model, text, n_samples=n_samples, temperature=temperature)

    # Compute the overlap scores
    scores = weighted_avg_overlap(samples, log_probs, similarity_fn)

    # Decode samples
    decoded_translations = [tokens_to_text(s, id_to_word).numpy().decode('utf-8') for s in samples]

    # Find the key with the highest score
    max_score_key = max(scores, key=lambda k: scores[k])

    # Get the translation
    translation = decoded_translations[max_score_key]

```

```

translation = decoded_translations[max_score_key]

```

```

english_sentence = "I love languages"

translation, candidates = mbr_decode(trained_translator, english_sentence, n_samples=10,
print("Translation candidates:")
for c in candidates:
    print(c)

```

```

Translation candidates:
eu adoro idiomas em frances .
eu adoro idiomas a vida .
eu gosto de idiomas que eu exagera a vida . eu estou .
eu adoro idiomas o dia .
eu amo idiomas ao olho ?
eu amo linguas os linguas .
eu adoro idiomas para idiomas .
eu adoro idiomas o professor .
eu adoro linguas rapidamente .
eu adoro idiomas para [UNK] .

```

Selected translation: eu adoro idiomas para idiomas .

Congratulations! Next week, you'll dive deeper into attention models and study the Transformer architecture. You will build another network but without the recurrent part. It will show that attention is all you need! It should be fun!

Keep up the good work!